

Chapter 1

Introduction

This dissertation describes the implementation and evaluation of an activity classifier using accelerometer data captured simultaneously from a smartphone and a smartwatch.

The classifier using data from both sources outperforms a classifier using only smartphone data, and the classifier that uses only smartphone data outperforms a classifier using only smartwatch data.

1.1 Motivation

Wearable devices are set to become the next big technology trend. Wrist-worn wearables, including smartwatches, formed the majority of the 21m wearable devices sold year. Analysts predict the Apple Watch will sell between 20m and 40m in its first nine months [10].

One of the primary appeals of wearables is their ability to sense. Like smartphones before them, smartwatches will enhance the ability to collect data about people. This data is important to consumers, who purchase specialised wearables to measure activity, sleep patterns and calorific intake. The data's research potential is also laudable — Apple's ResearchKit will allow medical researchers to access data about their patients with greater ease than ever before [7].

Accurate activity classification therefore has many academic and commercial applications. To be marketable, activity classification solutions must use current

consumer devices. Though rudimentary activity classification is available on Android smartphones, an approach that utilises simultaneous collection from a smartphone and smartwatch has not been investigated in any detail.

For that reason, this dissertation details the implementation of accelerometer data collection using current consumer devices (an Android smartphone and Android Wear smartwatch), classifies a user's activities and compares this classification accuracy to using only smartphone data and using only smartwatch data.

1.2 Challenges

This project requires knowledge of a variety of disparate areas in computer science.

Writing software for mobile devices requires knowledge of their paradigms and nuances. Mobile devices are also subject to battery life and computational power constraints and particular care must be taken to build a solution that works in practice. A project that utilises built-in sensors also requires an understanding of the features and limitations of those sensors and good knowledge in the APIs that are provided to access them.

The sensors also output data at a high rate and care must be taken to correctly handle the performance and concurrency issues that may arise. Storage and transfer of large amounts of raw data, especially on a memory-limited device such as a smartwatch, also requires special consideration.

The data processing aspects of the project will require an understanding of digital signal processing, Fourier methods, artificial intelligence and machine learning, and statistics.

1.3 Related Work

Activity classification using accelerometer data from body-mounted devices is an active area of research. I highlight three papers and discuss their similarity to this problem. Summaries of their work are found in Table 1.1.

Bao *et al.* [2] detect physical activities using five biaxial accelerometers worn on different parts of the body: hip, wrist, ankle, arm and thigh. They find that accuracy is not significantly reduced when using just thigh and wrist accelerometers. Furthermore, recognition rates for thigh and wrist data resulted in the highest recognition accuracy among all pairs of accelerometers, with over a 25% improvement over the best single accelerometer results. This supports the viability of this project, with the improvement of being able to use triaxial accelerometers found in consumer smartphones and smartwatches.

Long *et al.* [6] use a single triaxial accelerometer placed on the wrist and use it to achieve an 80% activity classification accuracy in five activities. However, only 50% of all cycling is correctly classified. Bao *et al.* achieve an accuracy of $> 92\%$ by using thigh and wrist data. This would suggest that wrist data alone is not sufficient to accurately classify certain types of activity. Cycling requires periodic leg motion (pedalling) while the hands and wrists move comparably little. Many of the features of motion used in activity classification require frequency domain analysis, and so data that contains periodic motion will be easier to recognise.

Atallah *et al.* [1] focus on two important facets of accelerometer-based activity classification: sensor location and useful features. Much like Bao *et al.* they use seven sensors on the chest, arm, wrist, waist, knee, ankle and ear. Of their analysed features, the averaged entropy over three axes, the mean of the pairwise cross-covariance of axes and the energy of a 0.2 Hz window around the main frequency divided by total energy are all highlighted as being highly ranked for distinguishing activities. However, this study neglects to use a decision tree classifier in its classification, recommended by both Bao *et al.* and Long *et al.*

	Bao <i>et al.</i> [2]	Long <i>et al.</i> [6]	Atallah <i>et al.</i> [1]
Activities	Walking, sitting & relaxing, standing, watching TV, running, stretching, scrubbing, folding laundry, brushing teeth, riding elevator, carrying items, computer work, eating or drinking, reading, bicycling, strength-training, vacuuming, lying down, climbing stairs, riding escalator	Walking, running, cycling, driving, sports	Lying down, preparing food, eating and drink- ing, socialising, reading, getting dressed, corridor walking, treadmill walking, vacuuming, wiping tables, corridor running, treadmill running, cycling, sitting down and getting up, lying down and getting up
Features	Mean, energy, correlation, entropy	Standard deviation, entropy, orientation vari- ation	Mean, variance, root mean square, entropy, correlation, range, energy, primary frequency, skewness, kurtosis
Classifiers	Decision table, nearest neighbour, decision tree, naive Bayes	Decision tree, principle compon- ent analysis, naive Bayes	K-nearest neigh- bors, naive Bayes
Overall accuracy	84%	80%	N/A

Table 1.1: Prior work on accelerometer-based activity classification

Chapter 2

Preparation

This chapter details the work done before the main implementation of the project was started. It details the devices chosen to implement this project and the reasons for choosing them. It then discusses the existing libraries and APIs available for those devices and for the required data processing. Finally, it describes software engineering techniques used.

2.1 Requirements analysis

The aim of the project is to classify activities based on accelerometer recordings from a consumer smartwatch and smartphone, and evaluate to what extent the smartwatch is better at helping to classify activities. The requirements to accomplish this can be split into two categories: data collection and data processing requirements.

Data collection requirements

1. access tri-axial readings from accelerometer on both the smartwatch and the smartphone;
2. store this accelerometer data temporarily on the internal memory of each device using suitable data structures;

3. transmit this data from the smartwatch to the smartphone using a suitable protocol;
4. store the data permanently on the smartphone, to enable transfer to the computer.

Data processing requirements

1. parse the data into a manipulatable format;
2. preprocess the data, including filtering and splitting into fixed-length bins;
3. extract features from each bin;
4. train classifier(s) on the extracted features;
5. test classifier and record evaluation statistics.

The remainder of this chapter describes work done to ensure these requirements could be fulfilled.

2.2 Introduction to signal processing

The output from any accelerometer is a time-series representing its acceleration. Effectively extracting information from this time-series is central to the success of this project. Knowledge of signal processing is therefore critical.

It is essential to capture as much of the movement as possible. Conversion from continuous physical acceleration to a discrete time-series requires sampling. The Nyquist-Shannon sampling theorem states that a signal can be exactly reconstructed from its samples if the sample rate is greater than twice the highest frequency of the signal.

The highest frequency of a physical activity is not well defined. The activities I hope to classify will vary in their periodicity. Some, like walking, will be very periodic, while others will have no period at all. Considering common period activities like cycling and walking, I anticipate that the frequencies that best describe movement will be present in the 0–5Hz range. Graphs of acceler-

ometer readings for each of the activities I attempt to classify are presented in section 3.2.

Frequency domain analysis

Much of the analysis of the accelerometer readings will be done in the frequency domain. A time domain signal can be converted into the frequency domain using a Fourier transform.

The discrete Fourier transform of a sequence of N complex numbers f_0, f_1, \dots, f_{N-1} is the sequence F_k , defined by:

$$F_k = \sum_{n=0}^{N-1} f_n \cdot e^{-2\pi i k n / N}$$

The power spectral density of a signal describes how power is distributed over different frequencies. One method of estimating the power spectral density is to take the square of the absolute value of the Fourier transform component:

$$PSD_k = \|F_k\|^2$$

Noise and filtering

The readings from the accelerometer are subject to noise, exhibited in figure 2.1, which plots readings from the x, y, and z axes during an hour long recording with the phone laying flat on a table.

Figure 2.2 plots the distribution of the magnitude of the acceleration, where the magnitude $\|x\| = \sqrt{x^2 + y^2 + z^2}$. The magnitude, which should be a constant $g \approx 9.81 \text{ m s}^{-2}$, is subject to normally distributed noise.

Figure 2.3 gives a normal probability plot of the same magnitude data. Points on a normal probability plot should form a straight line if they are normally distributed. The straight line of best fit exhibits a coefficient of determination, R^2 , which is very close to 1 and therefore it is very likely that the noise is normally distributed.




Figure 2.1: The x, y and z axis readings from an hour long accelerometer recording of the phone laying flat on a table. The readings contain noise.




Figure 2.2: Histogram of the magnitude $\|\mathbf{x}\| = \sqrt{x^2 + y^2 + z^2}$ of the data shown in Figure 2.1. The magnitude should measure $g = 9.81\text{ms}^{-2}$. The noise means the accelerometer data is imprecise. The mean of the data is less than g , which indicates the recording is also inaccurate.




Figure 2.3: A normal probability plot of the magnitude $\|\mathbf{x}\| = \sqrt{x^2 + y^2 + z^2}$ of the data shown in Figure 2.1. Data that is normally distributed will form a straight line when plotted in this way. This data is very likely to be normally distributed.

Noise can be reduced with the application of a low-pass filter. A low-pass filter attenuates signals with a higher frequency than some cutoff, such as the noise exhibited in the signal.

2.3 Hardware devices

The success of this project depends partly on correct selection and understanding of the devices used to collect data. The devices are both required to contain accelerometers accessible to developers.

Android devices were chosen as Android Wear was the most mature platform for developing with wearable devices at the time. It runs on the widest variety of devices and provides developer access to its sensors.

2.3.1 Smartphone

The smartphone chosen for development was the Google Nexus 5. Smartphone technology has advanced to the point that many Android smartphones are homogeneous with respect to this project — they all continue sufficient processing power, internal memory and an accelerometer capable of recording data.

The Nexus 5 contains a tri-axial accelerometer capable of recording measurements $\pm 2g$ on each axis, where $g \approx 9.81 \text{ m s}^{-2}$.

2.3.2 Smartwatch

The smartwatch chosen for development was the Samsung Galaxy Gear Live, running Android Wear. It pairs to any device running Android 4.4 or higher and communicates over Bluetooth.

Wearable devices not running Android typically run either Tizen, an open-source but not widely adopted operating system — such as the Samsung Galaxy Gear 2 — or a proprietary operating system that does not allow access to the raw accelerometer data, for example the Jawbone Up.

There is more differentiation in smartwatches as there is in smartphones, with them varying not just in screen size but also in screen format (round or rectangular), battery life, charging facilities and sensors. Table 2.1 presents an overview of possible smartwatch devices.

2.4 Libraries and APIs

This project makes use of existing libraries and APIs for the data collection, data handling and classification aspects of the project. Preliminary investigation into each of these areas was conducted.

2.4.1 Android Sensor API

The Android platform Sensor API is implemented using a publisher-subscriber model. Listeners to a particular sensor must be registered and must implement an `onSensorChanged()` method. The `onSensorChanged()` method is called whenever the sensor reports a new value. A `SensorEvent` object is provided, containing a timestamp at which the data was reported together with the new data.

The rate at which `onSensorChanged()` is called is user-‘suggested’; though it can be specified by the user, it can also be altered by the Android system. In practice, this means that the difference in timestamps not constant but is approximately equal to the specified delay. A histogram of timestamp differences for a particular 1 hour recording is given in figure 2.4.

Android provides both acceleration and linear acceration sensors, related by

$$\text{acceleration} = \text{linear acceleration} + \text{gravity}$$

They each provide a timestamp represented as a long and three float values representing the acceleration of each axis in m s^{-2} at that timestamp. Table 2.2 gives a graphical representation of the data returned.

Curiously, the timestamp returned as part of the data is documented only as “The time in nanosecond at which the event happened” [5]. Futher exploration reveals that the timestamp is not defined against any particular zero-base, but

Device	Samsung Galaxy Gear Live	Samsung Galaxy Gear 2	LG G Watch	Sony Smartwatch 3
Operating System	Android Wear	Tizen	Android Wear	Android Wear
Processor	1.2 GHz single-core Qualcomm Snapdragon 400	1.0 GHz dual-core Exynos 3250	1.2 GHz single-core Qualcomm Snapdragon 400	1.2 GHz quad-core ARM A7
Memory	512 MB RAM	512 MB RAM	512 MB RAM	512 MB RAM
Storage	4 GB	4 GB	4 GB	4 GB
Sensors	Touchscreen, Accelerometer, Gyroscope, Compass, Heart Rate Monitor	Touchscreen, Accelerometer, Gyroscope, Heart Rate Sensor, 2 MP Camera	Touchscreen, Accelerometer, Gyroscope, Compass	Touchscreen, Accelerometer, Gyroscope, Compass
Radios	Bluetooth 4.0 Low Energy	Bluetooth 4.0 Low Energy	Bluetooth 4.0 Low Energy	Bluetooth 4.0 Low Energy, GPS, NFC, Wi-Fi
Battery	300 mAh	300 mAh	400 mAh	420 mAh
Notes		Pairs only with Samsung devices		

Table 2.1: An overview of possible smartwatch devices. The Samsung Galaxy Gear Live was the device eventually chosen.




Figure 2.4: Histogram of the differences in successive timestamps of a one hour accelerometer recording from the Nexus 5 smartphone. The sample rate was set to 50 Hz. 0.02002s accounted for 75% of the differences. Thus the actual sample rate is approximately the user-suggested sample rate.

rather the time since the device was powered on [4, 8]. The implication of this for the project is that while the timestamp can be relied on for intervals between measurements, it cannot be used between different sets of recordings or across devices.

2.4.2 ES Sensor Manager

I explored this but didn't end up using it. Should I write about what it is and why I didn't end up using it?

Timestamp	X	Y	Z
ns	acceleration m s^{-2}	acceleration m s^{-2}	acceleration m s^{-2}
Long	Float	Float	Float
2 bytes	1 byte	1 byte	1 byte

Table 2.2: Data from the accelerometer sensor provided to the `onSensorChanged()` method.

2.4.3 Android Wear Data API

As discussed in section 2.3.2, the only radio present in the Samsung Galaxy Gear Live is Bluetooth. To transfer any recorded data from the watch, it must first be transferred to the paired smartphone. The Android Wearable Data Layer API allows communication between Android handheld and wearable devices. It provides three methods of communication between devices:

Data items provide data storage with automatic syncing;

Messages are good for remote procedure calls but do not carry data;

Asset objects for sending binary blobs of data.

The data layer synchronises data between the handheld and wearable. To do so, the Wearable Data Layer API requires the registration of a listener service, much like the Sensor API. The listener service listens for data layer events, such as the creation of asset objects or when messages are received.

2.5 Choice of tools

2.5.1 Programming languages

Java was chosen as it is the native programming language used on Android. Although it is possible to write code for Android in programming languages other than Java, for example by using the Java Native Interface, doing so would not benefit the project. Java is taught in Part 1A and Part 1B of the Computer Science Tripos. The Android SDK builds on principles covered in the course but

is complicated by having to manage interactions with the Android operating system.

XML is Android's standard markup language. All user-interface components are written in XML. The project includes a user interface to configure and control the recording of data.

Python 3.4 was chosen as the data processing language due to its easy of use and the strength of its data processing, signal processing and machine learning libraries:

NumPy a scientific computing library and the basis for the other three libraries below.

Pandas extensions to NumPy that enable easier processing of time-series data.

SciPy signal processing tools and other statistical features.

Scikit Learn machine learning classifiers and utilities to work with them.

All of NumPy, SciPy, Pandas and Scikit Learn are open-source and licensed under the BSD license.

2.5.2 Development Environment

Two powerful IDEs, Android Studio and PyCharm were used for the development of the Android app and the Python data pipeline respectively. Android Studio is available for free from Google, while PyCharm is provided free for educational use by JetBrains. Both include advanced debuggers.

Though the Android SDK contains a device emulator, it runs slowly and cannot simulate sensors. Developing the Android apps is therefore done by connecting them to a computer and running new versions of the code. This also enables access to the device's logs from the development environment. I made extensive use of logging to determine that the program was executing as expected.

2.6 Software engineering techniques

2.6.1 Development methodologies

I used a combination of development methodologies for the project. The data collection apps were developed using a waterfall methodology, while the data processing was developed using an Agile methodology.

Waterfall models are excellent when the end goals of the project are known and can be well specified. The goal of the data collection apps can be easily stated: write apps for the smartphone and smartwatch that will allow user collection of accelerometer data.

The data processing and machine learning elements of the project required an Agile methodology. The goal here is less well defined — classify activities with the greatest accuracy — and the implementation to achieve the goal is far more experimental.

2.6.2 Version control and backups

I used three separate Git repositories for the data collection code, the data processing code and the dissertation respectively. The Git repositories were synced to GitHub at each commit. Version control allowed me to follow a *implement–test–commit* pattern when writing code.

GitHub also served as one method of backup. Each GitHub repository is publicly accessible such that I can continue implementation even if my primary development computer crashed and I was also locked out of my GitHub account. In addition, I backed up periodically to Dropbox and to an external hard drive. The external hard drive backup retained old copies of files when they are updated. This gives four replications of my entire project, with two of these able to access previous versions of the code.

2.7 Summary

In this section I presented:

- an overview of digital signal processing;
- information on the smartphone and smartwatch used;
- details of key APIs used including the Android Sensor API and the Android Wear Data API;
- development tools and software engineering techniques.

Chapter 3

Implementation

3.1 Data collection

This section contains details of the components built to access the accelerometer data and transfer it to a computer.

Because both the smartwatch and the smartphone both run Android, it is possible to create components that are shared between the devices, reducing the amount of code I am required to write and to test, resulting in less redundancy, less complexity and ultimately a more reliable implementation. Both the `AccelerometerListenerService` and the `AccelerometerDataBlob` are shared between both devices.

3.1.1 Accessing the accelerometer

The `AccelerometerListenerService` is responsible for receiving readings from the accelerometer and delivering them to the data structure responsible for storage.

As described in Section 2.4.1, the Sensor API utilises a listener methodology. It is required to create and register a listener that implements `onSensorChanged()`.

Performance considerations

Because the accelerometer can update its values at a rate of over 50Hz, it is vital that any implementation of `onSensorChanged()` be non-blocking and ideally be very quick to execute. Any expensive computation or IO operation has to be moved to a separate thread.

If the execution of `onSensorChanged()` takes longer than $\frac{1}{\text{sample-rate}}$, requests for `onSensorChanged()` will queue and eventually lead to exhaustion of memory or dropping of data.

For this reason the data structure used, discussed in Section 3.1.2, is very light-weight and `onSensorChanged()` is only responsible for passing data to it.

Concurrency considerations

Because `onSensorChanged()` can be called at such a high rate, it is possible that new calls to the method can be made while previous calls are still executing. Data corruption could result from improper handling of asynchronicity.

The documentation for the Sensor API is not explicit about whether calls to `onSensorChanged()` queue on the same thread or whether they can be dispatched asynchronously. For this reason, the `AccelerometerListenerService` was designed to be thread-safe by using Java concurrency primitives.

Power consumption considerations

Recording data from the accelerometer can be computationally expensive. This increase in computational overhead translates to an increase in power consumption in battery powered devices such as the smartphone and the smartwatch. It is for this reason that care should be taken to minimise power usage where possible while still collecting all the required data.

One tradeoff that had to be made was between collection strategies. One strategy is to record data at a specified sample rate from when the recording is turned on until it is turned off. An alternative strategy is to record a window of data at set intervals and sleep the remainder of the time. For example, one might set the accelerometer to record 10 seconds of data every 50 seconds.

Though this strategy saves battery power as the device turns off the accelerometer between recordings, a continuous recording approach was taken in this project in order to have as much data as possible with which to train. In addition, the battery life was not severely impeded by the continuous recording approach.

Typically, Android will power off the display and later the CPU after a period of user-inactivity. Powering off the CPU means that the device will stop recording accelerometer data, and so it is required to maintain a wake-lock which keeps the CPU from powering off. It is also important to remember to release the wake-lock once accelerometer recording is complete. Otherwise, the device's CPU will remain on even when the device appears to be on standby, using battery.

Sampling rate

In ideal conditions, it would be sensible to sample at as fast a rate as possible: the resultant data can always be downsampled afterwards if it is not required. As per the Nyquist-Shannon sampling theory, discussed in Section 2.2, our sample rate should be greater than twice the highest frequency of the signal. Because it isn't possible to know what the highest frequency is going to be, it would be reasonable to sample at a far higher rate.

However, picking a very fast sample rate in this context has two potential downsides: battery life drain and the size of resultant data. I investigated whether either battery life or the size of the resulting data would be a limiting factor of sample rate.

The impact on power consumption of increasing the sample rate was negligible. One possible reason for this is that sampling using the accelerometer at all has high fixed costs and increasing the sample rate has lower marginal costs.

Recall from Table 2.2 that each measurement has a total size of 20 bytes. At a sample rate of 50Hz, we produce data at approximately 1 KBps or 3.6 MB per hour. The most memory-constrained device is the smartwatch, which only has 512 MB of RAM but 4 GB of internal storage. A data structure that stores the accelerometer data to the internal storage rather than to memory is required, but a sample rate of 50Hz produces a storable amount of data on reasonable-length activity recording.

Another potential concern regarding data size is the transfer from the smartwatch to the smartphone. The only connection available is Bluetooth. The Bluetooth connection empirically has a maximum transfer rate of no more than 150 Kbps, meaning an hour of activity data will take approximately 30 seconds to transfer.

3.1.2 Storing accelerometer data

The data structure to hold the accelerometer data is required to be:

- **fast** because it will be accessed many times per second and cannot block;
- **on-disk** rather than in-memory, because the smartwatch may not have enough free memory to store all the accelerometer data for lengthy recordings;
- **thread-safe** as it is unclear whether calls to `onSensorChanged()` are queued or concurrent.

The data structure decided on was a temporary random-access file with buffered writing. The data is written as bytes through an output buffer. The output buffer is maintained in memory and is flushed when it reaches capacity. The capacity of the output buffer was set to 20000 bytes as data is only written in multiples of 20 bytes and the smartwatch is comfortably able to keep 20 kb in memory. This equates to data being saved to disk approximately every 20 seconds.

3.1.3 Transmitting accelerometer data

The accelerometer data has to be transmitted from the smartwatch to the smartphone before it can be transferred to a computer. As discussed in Section 2.4.3, there are two methods to transfer data between the smartwatch and the smartphone: a `DataItem` and an `Asset`. Their advantages and disadvantages with respect to this project are highlighted in Table 3.1.

Because the `DataItem` has a 100 KB limit, an alternate transmission and storage system would have to have been built, where the smartwatch collects 100 KB

	DataItem	Asset
Advantages	<ul style="list-style-type: none"> • no separate data fetching step • simpler, more reliable receiver code • negligible transmission time 	<ul style="list-style-type: none"> • no hard size limit • can create an Asset from a File without storing it in memory
Disadvantages	<ul style="list-style-type: none"> • 100 KB size limit • have to insert byte arrays 	<ul style="list-style-type: none"> • some constructors don't seem to work • transmission of large files takes a noticeable amount of time

Table 3.1: Advantages and disadvantages of using the DataItem and Asset to transmit data from the smartwatch to the smartphone.

of data and sends that to the smartphone while it continues to record. It is then reassembled at the smartphone receiver.

I consider this solution inferior to the Asset implementation, which allows transmission of any size of data.

3.1.4 User interface

Smartphone

Smartwatch

3.2 Activity analysis

I then used the smartphone and smartwatch apps to collect data over a range of activities. This section provides the graphs of the magnitude and the Fourier transform of the magnitude. I used these graphs to better understand the accelerometer readings resulting from different activities and justify the utility of the features I hope to extract from the data.

For each activity, I graph an arbitrary ten second snippet as recorded by both the phone and watch. The ten second snippet has been low-pass filtered with a critical frequency of 5Hz, as described in Section 3.3.1. The Fourier transform for the same snippet is also provided, so as to display the recording in the frequency domain as well as the time domain.

For all recordings, the watch was worn tight on the non-dominant left hand and the phone was kept in the right side trouser pocket.

3.2.1 Computer use

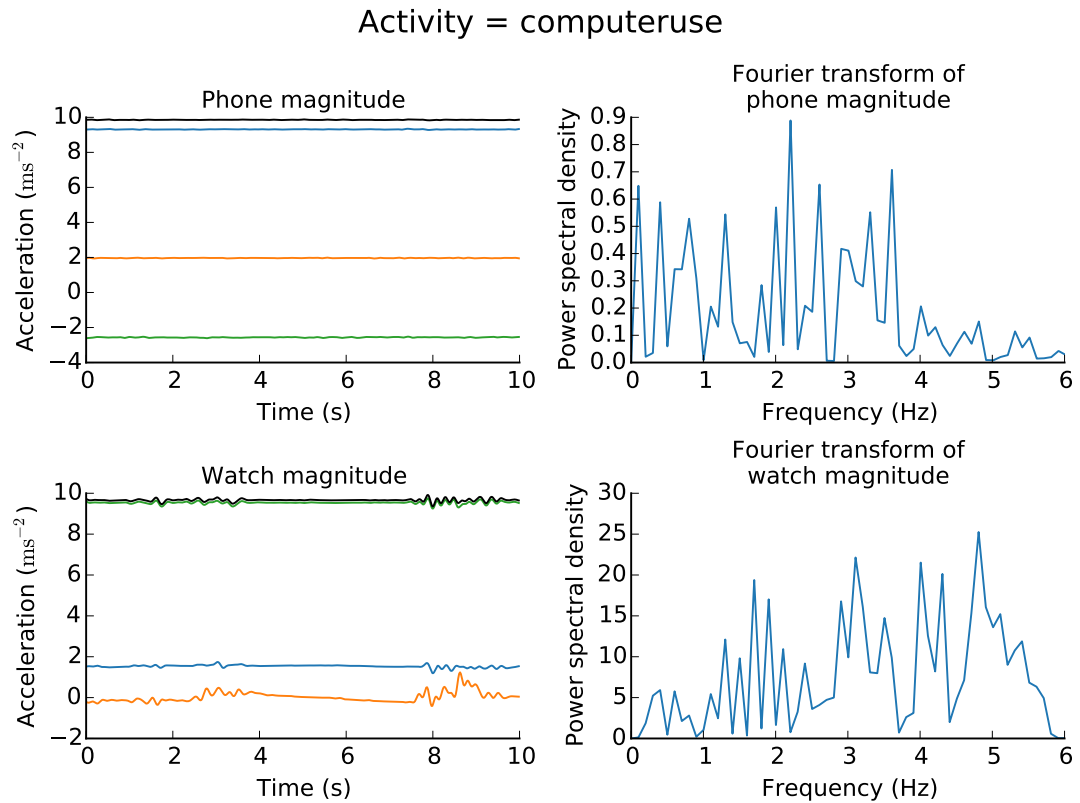


Figure 3.1: Ten seconds of phone and watch data from a computer use activity together with their Fourier transforms.

Computer use is predominantly typing or using a laptop trackpad while seated. There is very little movement in the phone, as the leg is mostly stationary. The

watch exhibits some periodic movement punctuated by periods of inactivity. This is presumably from typing a word and then pausing.

The Fourier transform of the watch magnitude indicates that watch movement is also aperiodic, as one might expect from typing with frequent pauses.

3.2.2 Walking

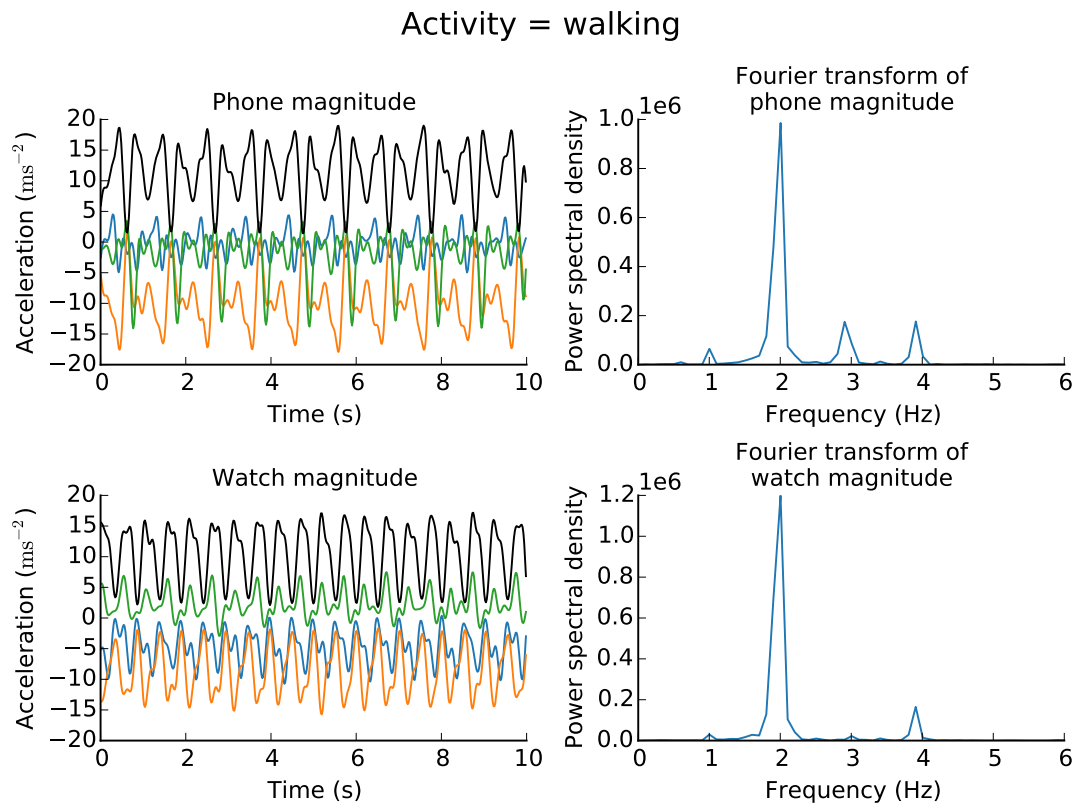


Figure 3.2: Ten seconds of phone and watch data from a walking activity together with their Fourier transforms.

Walking is among the most periodic of all the activities investigated. There is a very strong peak at about 2 Hz, a typical walking pace. There are smaller local maxima at 1, 3 and 4 Hz in the Fourier transform of the phone magnitude, and peaks at 2 Hz and 4 Hz in the watch magnitude. The rest of the Fourier transform is relatively flat.

A good distinguishing feature for between activities that exhibit strong periodicity is the spectral flatness measure, discussed in Section 3.3.2.

3.2.3 Cycling

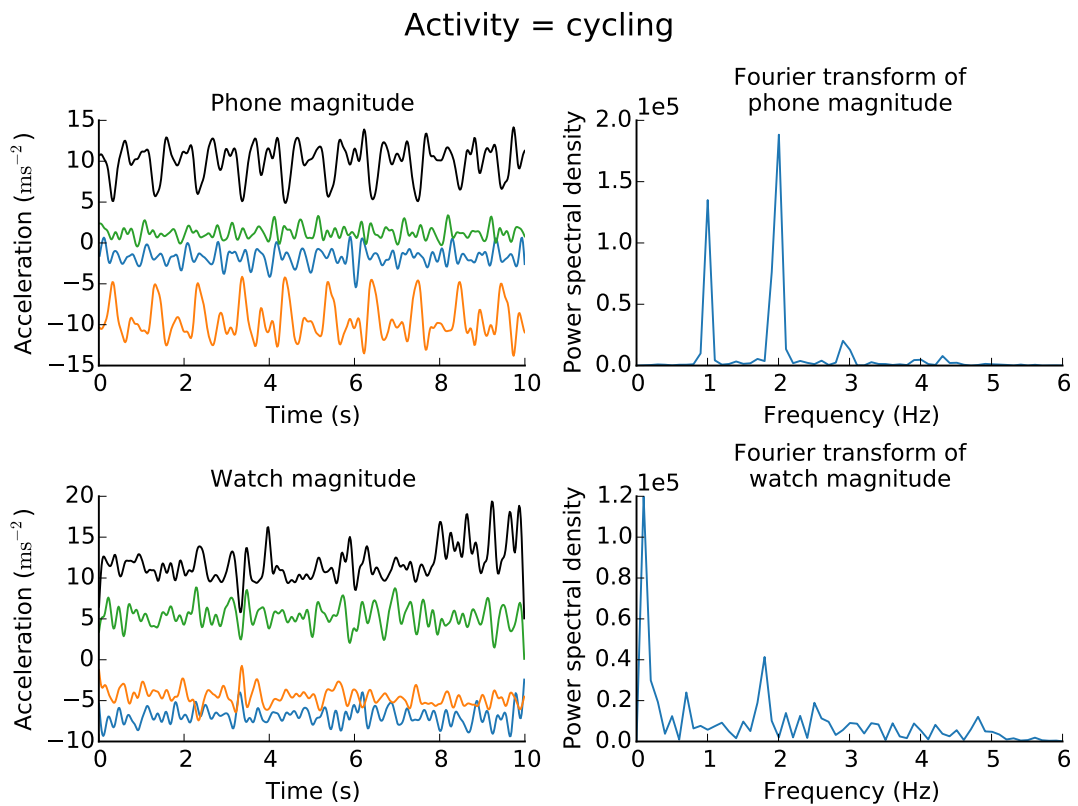


Figure 3.3: Ten seconds of phone and watch data from a cycling activity together with their Fourier transforms.

Cycling is another activity with high periodicity in the phone measurement, but, unlike walking, does not have much periodicity in the watch measurement. This is primarily because of the changing position on the handlebars.

3.2.4 Standing

Standing as an activity exists to test differentiation between other upright activities such as toothbrushing and fussball. The phone is largely stationary while the watch moves between certain key standing positions (e.g. arm hanging, in pocket, on hip etc.).

The mean and the standard deviations of the x, y, and z axes of the phone will show standing activities. The lack of movement in the watch could

3.2.5 Playing fussball

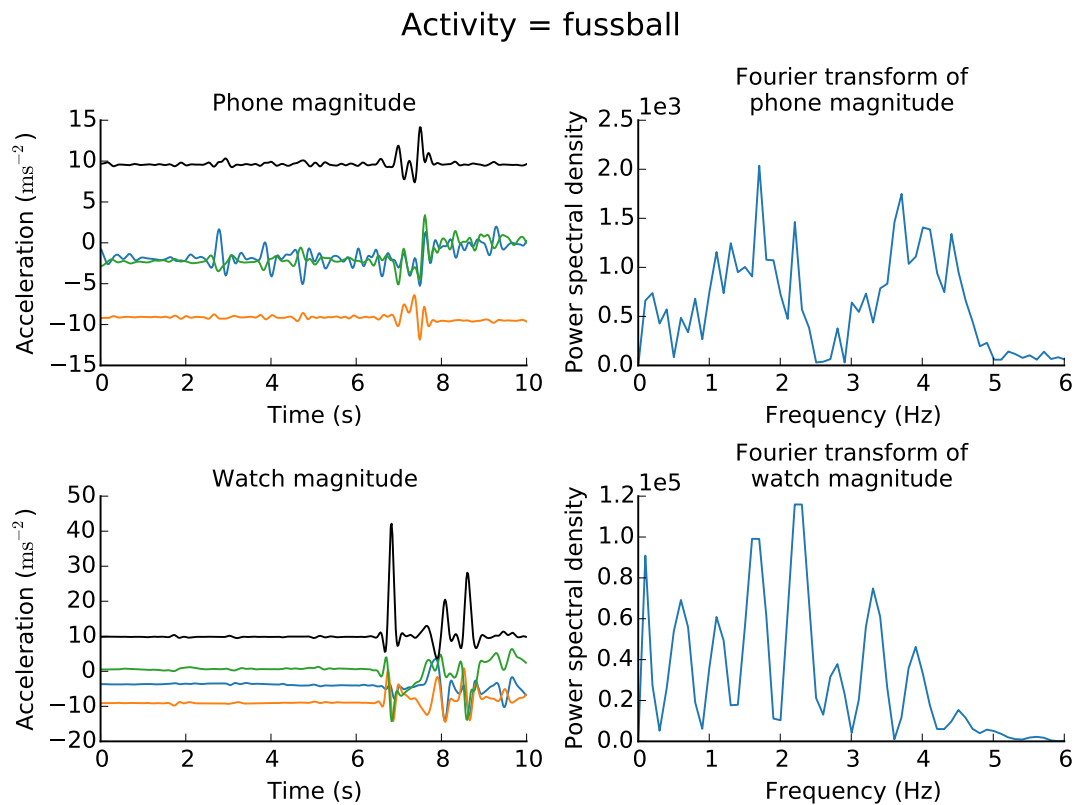


Figure 3.4: Ten seconds of phone and watch data from a fussball activity together with their Fourier transforms.

Fussball (also known as table football) is characterised by periods of inactivity

followed by sharp acceleration in the watch measurement. The maximum of the magnitude should in theory differentiate this activity from others.

However, the maximum is by definition a statistic that is highly sensitive to outliers and so other activities may exhibit a similarly high magnitude even if it is not characteristic of that activity. A better metric might be a count of the number of data points that exceed a certain threshold magnitude e.g. 40m s^{-2}

3.2.6 Stair climbing

Activity = stairs

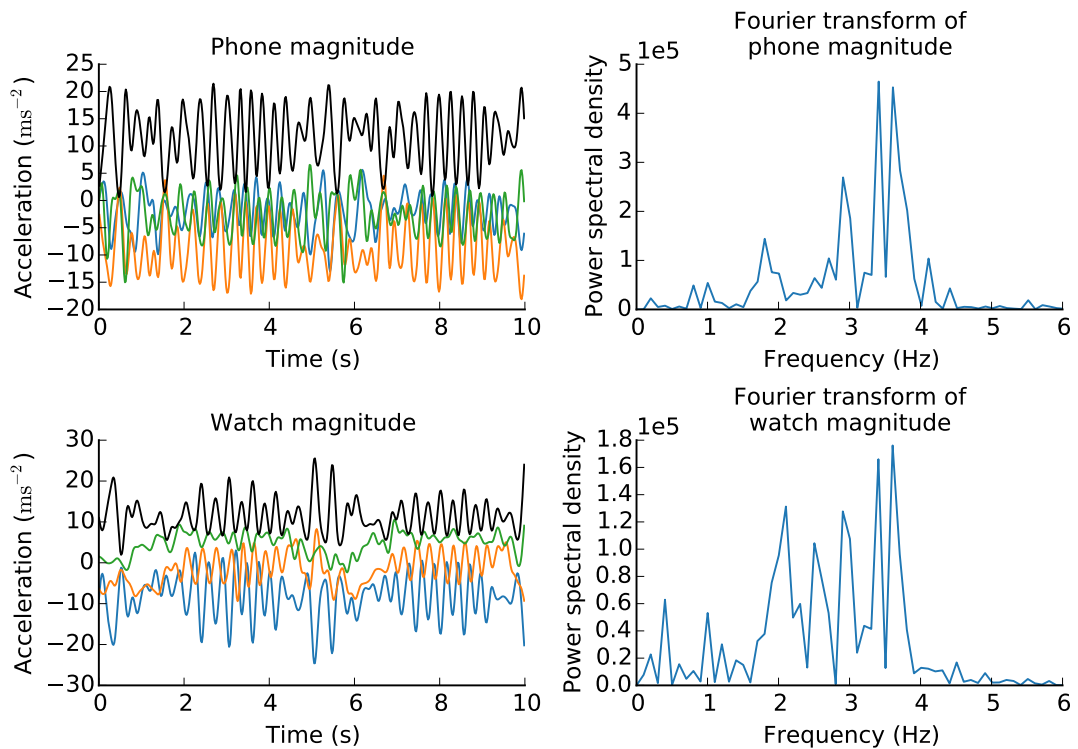


Figure 3.5: Ten seconds of phone and watch data from a stair climbing activity together with their Fourier transforms.

Stairs were climbed in the computer lab. A single recording contains climbing and descending stairs. Stairs were climbed either one at a time or two at a time, with the hand either loose or holding the handrail.

This means that stair climbing exhibits a variety of frequencies — descending stairs is typically done at a much quicker pace than climbing stairs, for example. This particular example seems to be drawn from stair descending.

Stairs can be differentiated from walking by noting that though there are still definite periods apparent in the Fourier transform, the peaks are not quite as clear as they are while walking.

3.3 Data processing

The data processing pipeline was written in Python, because of the strength of its numerical, statistical and machine learning libraries. The data processing was done on a computer as opposed to directly on the phone because of the computational demand required of training a classifier. In a real-world use-case, one could imagine classification occurring on a cloud server.

3.3.1 Importing and preprocessing

Import

Each recording is stored in a separate binary file. The filename is always of the form: <timestamp>-<recorder>-<device>-<activity>.dat.

NumPy includes methods to specify the types of binary data in a file and create an array from it. These are used to great effect to convert the binary data back into longs and floats.

Access to the data files is then via a SQLite database, using the timestamp as the unique ID. The database allows easier access to individual records and, for example, all recordings of a certain activity.

Preprocessing

Data is preprocessed before feature extraction.

The first step is to drop the first and last 10 seconds of each data recording. The justification for this is that these accelerometer recordings will not actually be representative of the activity we are attempting to classify. Rather, they will primarily be recording the starting and ending of an activity.

For each data recording, the magnitude of the acceleration $\|\mathbf{x}\| = \sqrt{x^2 + y^2 + z^2}$ was calculated. Patterns in the magnitude were found to be more distinguishing than any of the features extracted from the three axes individually. The magnitude is orientation-invariant, which gives better results when considering that a wrist may move in the same way but may be oriented slightly differently. In this scenario, periodicity will still be observed in the magnitude but may not be observed in each of the three axes individually.

Data is then filtered. As discussed in Section 2.2, the data recorded by the accelerometer is subject to noise. Reducing the effect of this noise will produce a signal in which it is easier to observe the underlying patterns produced by movement.

A fifth-order Butterworth Filter with a critical frequency of 5 Hz was used. The Butterworth Filter was chosen because it has no gain ripple in the pass band or the stop band and because the slow cutoff is not a problem for our application, as the frequencies of activities we are concerned with are far less than the frequency of the noise. A graph of the frequency responses for several Butterworth Filters is given in Figure 3.6.

Windows

Each data recording is split into 10 second windows. Features are extracted from each of these bins individually. Splitting into windows allows the production of multiple feature rows from the same recording. In theory, every window for a particular activity should exhibit extracted features that are consistent. A 10 second window was picked as a balance between producing enough feature rows from collected data and ensuring that several cycles of an activity were included.




Figure 3.6: Frequency response for Butterworth filters of different orders. Each has a critical frequency of 5 Hz. A fifth-order Butterworth filter was used to filter the noise from the accelerometer data.

No of features	Description of each feature
4	Mean of each axis and magnitude
4	Standard deviation of each axis and magnitude
4	Maximum amplitude of each axis and magnitude
4	Median absolute deviation of each axis and magnitude
3	Pairwise correlation coefficient of each of the three axes
1	Spectral flatness of magnitude
1	Spectral entropy of magnitude
1	Frequency of maximum amplitude in the power spectrum of the magnitude

Table 3.2: A summary of extracted features.

3.3.2 Feature extraction

In total, 22 features are extracted from each window. The smartwatch and smartphone data are treated as separate windows for the purposes of feature extraction. A summary of the extracted features is given in Table 3.2. This section goes on to describe and justify each of these features.

Mean of each axis and magnitude

The arithmetic mean \bar{X} defined as

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n x_i$$

was calculated for each of the x, y and z axes and also for the magnitude.

The arithmetic mean does not encode a lot of data, but is useful for determining primary orientation during the activity. For example, computer use has a z mean which is close to gravitational acceleration while the others are near zero. This indicates the devices primarily points upward during this activity.

Standard deviation of each axis and magnitude

The standard deviation σ defined as

$$\sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (\bar{x} - x_i)^2}$$

was calculated for each of the x, y and z axes and also for the magnitude.

Maximum amplitude of each axis and magnitude

The maximum x_{\max} defined as

$$x_{\max} = \max_i x_i$$

Median average deviation of each axis and magnitude

The median average deviation x_{mad} defined as

$$x_{\text{mad}} = \text{median}_i (|x_i - \text{median}_j(x_j)|)$$

Though this measure follows the same trend as the standard deviation, its use of the median ensures it is a robust statistic i.e. one that is resistant to outliers and offers good performance on data that is not normally distributed. This is a desirable characteristic in activities such as futsal, as recorded from the wrist.

Pairwise correlation coefficient of each of the three axes

The covariance is a measure of how much two random variables change together. The covariance of two random variables X and Y , $\text{Cov}(X, Y)$ is defined as

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \bar{X})(Y - \bar{Y})]$$

The correlation coefficient, $\text{Cor}(X, Y)$, of two random variables X and Y is the normalised covariance of the two random variables.

$$\text{Cor}(X, Y) = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$$

$\text{Cor}(X, Y)$ has a value between -1 and 1 , with 1 representing total positive correlation, 0 representing no correlation and -1 representing total negative correlation.

The correlation coefficient was calculated for each pair of axes, producing three features: $\text{Cor}(X, Y)$, $\text{Cor}(X, Z)$, and $\text{Cor}(Y, Z)$. The correlation coefficients give a measure of how much the axes move together during the recording. This encodes some information about the direction of movement.

Spectral flatness of magnitude

Spectral flatness is also known as the tonality coefficient. It is a measure of how noise-like or tone-like a signal is. White noise has a spectral flatness approaching 1 , while a pure tone has a spectral flatness approaching zero.

Spectral flatness is calculated from the power spectrum of the signal. Recall from Section 2.2 that the power spectrum is the squared magnitude of the Fourier transform of the signal.

If x_i represents the magnitude in the power spectrum of bin i , then the Flatness of a power spectrum $X = [x_1, x_2, \dots, x_N]$ is defined as

$$\text{Flatness}(X) = \frac{\sqrt[N]{\prod_{i=1}^N x_i}}{\frac{1}{N} \sum_{i=1}^N x_i}$$

The geometric mean can be expressed as a summation of logarithms rather than a product, giving an alternative formula for the Flatness that does not require a large product or an n th root. As our Fourier transforms are likely to be over 10000 elements long, avoiding the large product or the expensive n th root calculation is desirable. Following conversion to a logarithmic summation, the Flatness is:

$$\text{Flatness}(X) = \frac{\exp\left(\frac{1}{N} \sum_{i=1}^N \ln x_i\right)}{\frac{1}{N} \sum_{i=1}^N x_i}$$

Spectral flatness gives a measure of how periodic a signal is. Activities where there is very high spectral flatness, akin to white noise, are aperiodic. For example, futsal and standing have no associated period, while walking when measured through the smartphone has a clear period.

Spectral entropy of the magnitude

The spectral entropy of a signal is calculated as the entropy of its power spectrum. It is defined as:

$$\text{entropy} = - \sum_{i=1}^N x_i \log_2 x_i$$

Frequency of maximum amplitude in the power spectrum of the magnitude

The power spectrum of the magnitude shows at which frequencies the power in the signal is distributed. It is defined as the squared magnitude of the Fourier transform of the signal.

Rather than take the Fourier transform of the original signal, the mean of the signal was first subtracted. If the original signal oscillates about some non-zero offset, the Fourier transform will have a spike at the origin (0 Hz, or the DC component). To avoid incorrectly classing 0 Hz as the maximum, we subtract the mean from the original signal. The frequency of maximum amplitude of a power spectrum X is then:

$$f_{\max} = \operatorname{argmax}_f X_f$$

The frequency of maximum amplitude gives the principle frequency of the accelerometer data. For example, humans take between one and two steps in a second, and so we should expect a frequency of maximum amplitude in the 1 ~ 2Hz range. Indeed, this is empirically what we see.

3.3.3 Machine learning

Classification is supervised learning problem. This requires a classifier to be supplied with a set of instances, comprising sets of features, and a set containing a label for each instance. I use the following notation:

- a set of instances $X = \{X_1, X_2, \dots, X_N\}$ where each X_i is a vector of j features;
- a multiset of labels y where $y_i \in \{1, 2, \dots, K\}$ is the label for instance X_i ;

This set of labelled instances is referred to as the *training set*, which a classifier uses to generate a model. The classifier is then given a set of unlabelled instances, referred to as the *test set*. The true labels of the test set are kept, but unknown to the classifier. The classifier generates a prediction for the test set, and comparison of the predicted labels with the true labels from the basis of evaluation of the classifier.

I used four different classifiers:

- Proportionally stratified random classifier
- Naive Bayes classifier
- Decision Tree classifier
- Random Forest classifier

Though each of these classifiers is included in Scikit Learn, the mechanisms of each of these classifiers are explained.

Proportionally stratified random classifier

The proportionally stratified random classifier acts as a dummy classifier. It assigns labels to instances randomly, based on the proportion of the labels in the training set. This classifier is used to establish a baseline for evaluation.

Thus, this classifier completely ignores all feature information. Given the training set, the classifier generates, for each $k \in \{1, 2, \dots, K\}$ a count c_k of the number of times label k appears in the set of labels y i.e. the number of instances that are labelled k .

Then during the test phase, the outputted label is in each case randomly selected. The probability of a label k being output is c_k/N , where N is the total number of instances in the training set.

This classifier is only used to provide a baseline measurement against which we can compare other classifiers that do make use of the extracted features.

Naive Bayes classifier

The Naive Bayes classifier makes the naive assumption that all the features in the instance are independent. It then uses Bayesian probability to calculate the most probable class given the instance.

Mathematically we want to find $\mathbb{P}(y_k \mid \mathbf{X}_i)$ for each label y_k and each instance \mathbf{X}_i .

$$\begin{aligned}
 \mathbb{P}(y_k \mid \mathbf{X}_i) &= \mathbb{P}(y_k \mid x_1, x_2, \dots, x_j) && \mathbf{X}_i \text{ is a vector of features} \\
 &= \frac{\mathbb{P}(y_k)\mathbb{P}(x_1, x_2, \dots, x_j \mid y_k)}{\mathbb{P}(\mathbf{X}_i)} && \text{Bayes' rule} \\
 &\propto \mathbb{P}(y_k)\mathbb{P}(x_1, x_2, \dots, x_j \mid y_k) && \mathbb{P}(\mathbf{X}_i) \text{ is constant with respect to the label} \\
 &= \mathbb{P}(y_k)\mathbb{P}(x_1 \mid y_k)\mathbb{P}(x_2 \mid y_k) \cdots \mathbb{P}(x_j \mid y_k) && \text{Naive independence assumption} \\
 &= \mathbb{P}(y_k) \prod_{i=1}^j \mathbb{P}(x_i \mid y_k)
 \end{aligned}$$

The task then is to find the most probable label given the data:

$$k = \operatorname{argmax}_k \mathbb{P}(y_k) \prod_{i=1}^j \mathbb{P}(x_i \mid y_k)$$

To calculate the probability of a particular continuous feature value x_i given a label y_k , one can assume that the feature values follow a Gaussian distribution. Then, one can calculate the mean μ and the variance σ^2 for the value of the feature for a particular class.

During the test phase, one can calculate the probability that x_i takes its actual value v given each of the classes using the equation for a normal distribution parameterised by μ and σ^2 for that particular class:

$$\mathbb{P}(x_i = v \mid y_k) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(v - \mu)^2}{2\sigma^2}\right)$$

Decision Tree classifier

A decision tree classifier forms a tree where each node is a decision about a feature and labels appear as leaves.

The training procedure for a decision tree classifier builds the tree recursively. For a node m , let Q be the set of instance-label pairs in m . Create a possible split $\theta = (j, t)$ where j is a feature and t is some threshold value. Generate two subsets Q_{left} and Q_{right} , where

$$Q_{\text{left}}(\theta) = \{(\mathbf{X}_i, y_i) \mid x_j \leq t\}$$

$$Q_{\text{right}}(\theta) = Q \setminus Q_{\text{left}} = \{(\mathbf{X}_i, y_i) \mid x_j > t\}$$

The impurity G at m is a measure of how far from a perfect dichotomy the split θ is for Q . The impurity G requires an impurity metric H .

$$G(Q, \theta) = \frac{|Q_{\text{left}}(\theta)|}{|Q|} H(Q_{\text{left}}(\theta)) + \frac{|Q_{\text{right}}(\theta)|}{|Q|} H(Q_{\text{right}}(\theta))$$

There are two typical functions for the impurity metric H : the Gini impurity and the information gain.

The Gini impurity is given by:

$$H(\mathbf{X}) = \sum_{i=1}^K f_i(1 - f_i) = 1 - \sum_{i=1}^K f_i^2$$

The information gain impurity is given by:

$$H(\mathbf{X}) = \sum_{i=1}^K f_i \log f_i$$

where f_i is the proportion of instances labelled y_i in X . Note that both the Gini impurity and the information gain impurity reach their minimum value, 0, when X contains only a single class.

The task then for the training phase of the decision tree classifier is to select the split θ^* that minimises the impurity:

$$\theta^* = \underset{\theta}{\operatorname{argmin}} G(Q, \theta)$$

The classifier then recursively repeats this process for Q_{left} and Q_{right} unless:

- Q_{left} or Q_{right} contain a single class, at which point they become a leaf node;
or
- a pre-specified maximum depth has been reached; or
- the number of samples sent to the child node is less than a pre-specified minimum.

The introduction of a maximum depth and minimum sample size attempts to reduce the tendency for decision tree classifiers to overfit to the training data.

The Scikit-Learn implementation of decision tree classifiers uses an optimised version of the CART (Classification and Regression Trees), developed by Breiman *et al.*[3].

A single CART model is easy to interpret, as it can be illustrated as a set of binary decisions.

Random Forest classifier

The random forest classifier is an example of an ensemble classifier, one that makes use of a set of other classifiers. The random forest classifier trains a collection of decision tree classifiers. Each decision tree classifier finds the best split from a random subset of features, rather than the absolute best split. The outputs from all the decision tree classifiers are then aggregated by the random forest classifier by outputting the modal label.

Decision tree classifiers are prone to overfitting, especially when they are deep. The random forest classifier reduces variance by excluding different features from the training set, thus is less prone to overfitting.

The cost of a random forest classifier over one single decision tree classifier is the increase in training time and memory requirements and also the loss in interpretability of the resulting tree structure. The increased computational time required of a random forest classifier is not such a detriment in this offline processing task, but one might choose either a single decision tree or even a naive Bayes classifier if attempting to classify activities live.

3.4 Summary

Chapter 4

Evaluation

4.1 Data collection method and data collected

4.2 Evaluation process

Once the data was gathered, it was processed and features were extracted according to Section 3.3.

The data set was then split into training set and a testing set. This was done with a stratified shuffling splitter. This means that the proportional distribution of true labels in the testing set follows that of the training set. Whether an instance is placed in the training set or the testing set is chosen at random.

The split was performed 10 times, with elements chosen at random each time. The splitter does not guarantee that the same split will not be made on subsequent splits, but such an event is unlikely given the size of the dataset.

The data set was split in the ratio 50:50 training:testing.

Once the data had been split, the labels were stripped from the testing data and set aside. Then, three separate instances of each of the four classifiers was trained on the training set. Each of the three instances only had access to features extracted from the phone accelerometer data, features extracted from the watch accelerometer data and both the phone-extracted and the watch-extracted features respectively.

Each of the three instances were then tested on the test set and their evaluation metrics — confusion matrix and F_1 measure — were calculated by comparing the true labels to the classified labels.

4.3 Evaluation metrics

Two primary methods of evaluation are used in this project: F_1 measure and the confusion matrix. This section details these methods of evaluation. These metrics are discussed in Sokolova *et al.*[9]

4.3.1 F1 measure

In order to define the F_1 measure, it is necessary to first define precision and recall. In a

Precision , defined as $\frac{\text{True positives}}{\text{True positives} + \text{False positives}}$, is the proportion of instances of a particular label in which the classifier's labels agree with the true labels.

Recall , defined as $\frac{\text{True positives}}{\text{True positives} + \text{False negatives}}$, is the proportion of all the instances with a particular label which are labelled as such.

The F_1 measure is then defined as the harmonic mean of precision and recall:

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

We use the F_1 measure rather than precision or recall in isolation because neither provides sufficient information. It is trivial to maximise recall in isolation: simply label everything. A precision of 1 indicates everything we have returned is correct, but does not take into account how many instances we have missed.

The F_1 measure reaches its best value, 1, when both precision and recall equal 1. That is when all the classifier's labels agree with all the true labels and none of the true instances have been mislabeled.

Precision and recall, and thus the F_1 measure, are only defined for a single class (i.e. a binary classification problem). As this is a multi-class classification problem, the F_1 measure is reported for each class separately.

Because multiple F_1 measures are calculated through separate splits, the mean of the F_1 measure is given together with its standard error. The standard error is given by:

$$SE_{F_1} = \frac{\sigma}{\sqrt{n}}$$

where σ is the standard deviation of all measurements of the F_1 measure and n is the number of trials.

4.3.2 Confusion matrix

A confusion matrix lists the true labels as rows and the classified labels as columns. Any given cell c_{ij} is a count of the number of instances with a true label of i that were classified as label j . Confusion matrices display all classifications and misclassifications.

The confusion matrices presented in this project are the additions of confusion matrices from separate trials.

4.4 Phone-only measurements

4.5 Watch-only measurements

4.6 Phone and watch measurements

4.7 Comparison

4.8 Feature importances

Bibliography

- [1] Louis Atallah et al. "Sensor placement for activity detection using wearable accelerometers". In: *Body Sensor Networks (BSN), 2010 International Conference on*. IEEE. 2010, pp. 24–29.
- [2] Ling Bao and Stephen S Intille. "Activity recognition from user-annotated acceleration data". In: *Pervasive computing*. Springer, 2004, pp. 1–17.
- [3] Leo Breiman et al. *Classification and regression trees*. CRC press, 1984.
- [4] *Documentation Enhancement: SensorEvent timestamp*. 26th Apr. 2013. URL: <https://code.google.com/p/android/issues/detail?id=7981> (visited on 28/03/2015).
- [5] Google. *SensorEvent — Android Developers*. URL: <http://developer.android.com/reference/android/hardware/SensorEvent.html> (visited on 28/03/2015).
- [6] Xi Long, Bin Yin and Ronald M Aarts. "Single-accelerometer-based daily physical activity classification". In: *Engineering in Medicine and Biology Society, 2009. EMBC 2009. Annual International Conference of the IEEE*. IEEE. 2009, pp. 6107–6110.
- [7] *ResearchKit for Developers*. 23rd Mar. 2015. URL: <https://developer.apple.com/researchkit/>.
- [8] *SensorEvent timestamp field incorrectly populated on Nexus 4 devices*. 13th June 2013. URL: <https://code.google.com/p/android/issues/detail?id=56561> (visited on 28/03/2015).
- [9] Marina Sokolova and Guy Lapalme. "A systematic analysis of performance measures for classification tasks". In: *Information Processing & Management* 45.4 (2009), pp. 427–437.

- [10] “Wearable technology: The wear, why and how”. In: *The Economist* (14th Mar. 2015). URL: <http://www.economist.com/news/business/21646225-smartwatches-and-other-wearable-devices-become-mainstream-products-will-take-more>.